

The Hidden Energy Crisis in Computing

A Whitepaper on Software's Environmental Impact

David Jean Charlot, PhD

Open Interface Engineering, Inc. (openIE)

david@openie.dev

February 2026 | Version 1.0

Executive Summary

The global computing infrastructure (data centers, personal devices, network equipment, and embedded systems) now consumes approximately **4-5% of global electricity production** [1]. This figure is growing at 7-10% annually, outpacing efficiency improvements in hardware. While attention has focused on data center design and renewable energy procurement, a critical factor remains largely unaddressed: **the energy efficiency of software itself**. This whitepaper examines how programming language design and software architecture decisions contribute to energy consumption, and outlines pathways toward more sustainable computing practices.

1. The Scale of Computing's Energy Footprint

1.1 Data Centers

Global data centers consumed an estimated **240-340 TWh** of electricity in 2022, roughly equivalent to the entire energy consumption of the United Kingdom [1]. The International Energy Agency projects this could reach **1,000 TWh by 2026**, driven primarily by AI workloads [2].

Key contributors: - **Compute workloads**: 40-50% of data center energy - **Cooling systems**: 30-40% of data center energy - **Storage and networking**: 10-20% of data center energy

1.2 End-User Devices

The billions of smartphones, laptops, tablets, and IoT devices worldwide collectively consume substantial energy [3]:

- **6.8+ billion smartphones** running applications continuously
- **2 billion personal computers** executing desktop software
- **15+ billion IoT devices** performing embedded computations

Each inefficient app, each poorly optimized algorithm, each wasteful background process contributes to aggregate energy demand.

1.3 Artificial Intelligence

AI workloads represent the fastest-growing segment of computational energy demand [4]:

- Training GPT-3 consumed approximately **1,287 MWh** of electricity
 - Training a large language model can emit **300+ tonnes of CO₂**
 - Inference at scale multiplies training costs across millions of queries
 - Global AI energy consumption may reach **134 TWh by 2027**, comparable to the Netherlands [2]
-

2. The Software Factor

2.1 Why Software Matters

Hardware efficiency improvements have historically followed Moore's Law: approximately 2x every 18-24 months. However, software inefficiency often negates these gains through what's known as Wirth's Law: "Software is getting slower more rapidly than hardware is getting faster" [5].

- A 2x hardware improvement provides no benefit if software becomes 2x less efficient
- Software complexity tends to grow faster than hardware speed
- New abstraction layers add overhead without adding proportional value

2.2 The Language Layer

Programming languages are the foundation of all software. Their design decisions ripple through every program written in them. A landmark 2017 study measured the energy consumption of 27 programming languages across 10 benchmark problems [6]:

Memory Management - Garbage-collected languages (Java, Python, JavaScript, Go) incur energy overhead for runtime memory management - Collection cycles cause CPU activity spikes and prevent optimal power states - Studies show GC overhead ranges from 5–25% of total program energy consumption, varying by workload intensity and heap pressure [7]

Runtime Systems - Interpreted languages execute 10-100x more instructions than compiled equivalents - JIT compilation provides partial mitigation but adds its own overhead - Python consumes **75.88x more energy** than C for equivalent computations [6]

Abstraction Overhead - High-level abstractions often compile to inefficient instruction sequences -

Virtual function calls prevent optimization - Generic programming can lead to code bloat

2.3 Measurement Challenges

A fundamental problem: most developers have no visibility into their software's energy consumption.

- Profilers measure time, not energy
- Cloud bills show cost, not consumption
- No feedback loop exists for energy optimization

Without measurement, improvement is impossible.

3. The Path Forward

3.1 Energy-Aware Language Design

New programming languages can address energy efficiency from first principles:

Deterministic Resource Management - Ownership and borrowing systems eliminate garbage collection - Memory is freed at predictable points, enabling optimal power management - No runtime overhead for memory safety - Languages like Rust demonstrate this approach achieves energy efficiency comparable to C [6]

Energy Annotations - Allow developers to specify energy budgets for code sections - Compiler estimates and checks budgets at compile time - Energy becomes a first-class concern, not an after-thought

Efficient Default Choices - Value semantics by default (avoid heap allocation overhead) - Compile-time computation where possible - Direct hardware access without abstraction penalties

3.2 Heterogeneous Computing

Modern processors include specialized units for different workloads:

- **GPUs**: Parallel numerical computation (up to 10x more energy-efficient for suitable workloads) [\[8\]](#)
- **TPUs/NPUs**: Machine learning inference (up to 30–80x more efficient than CPUs for ML) [\[9\]](#)
- **DSPs**: Signal processing
- **Secure enclaves**: Cryptographic operations

Each specialized unit is dramatically more energy-efficient for its intended workload than a general-purpose CPU. Languages that make heterogeneous computing accessible can achieve order-of-magnitude energy savings.

3.3 Compiler Optimization for Energy

Traditional compilers optimize for speed. Energy-aware compilers can target different metrics [\[10\]](#):

- Minimize memory traffic (DRAM access consumes **~100x more energy** than cache access)
- Prefer SIMD operations (more work per instruction fetch)
- Enable processor power states (avoid blocking busy-waits)
- Reduce code size (instruction cache efficiency)

3.4 Measurement and Feedback

Making energy visible enables optimization:

- Hardware power monitoring integration (Intel RAPL, ARM Energy Probe) [\[11\]](#)
- Energy profiling tools
- Continuous integration energy tracking
- Energy regression detection

4. Empirical Evidence

4.1 Programming Language Energy Consumption

The following data is from Pereira et al.'s comprehensive study “Energy Efficiency across Programming Languages” [\[6\]](#):

Language	Relative Energy	Relative Time	Relative Memory
C	1.00	1.00	1.00
Rust	1.03	1.04	1.03
C++	1.34	1.56	1.34
Java	1.98	1.89	6.01
Go	3.23	2.83	1.05
JavaScript	4.45	6.52	4.59
TypeScript	21.50	46.20	4.69
Python	75.88	71.90	2.80

4.2 Data Structure Selection

Data structure choice affects memory access patterns and cache behavior:

Structure	Access Pattern	Cache Efficiency
Linked List	Random	Poor
Array	Sequential	Excellent
Hash Table	Random	Moderate
B-Tree	Localized	Good

Cache-efficient data structures can reduce energy consumption by 50% or more for memory-bound workloads [\[12\]](#).

4.3 The Business Case: Total Cost of Ownership

Energy efficiency translates directly to cost savings. For organizations operating at scale:

Metric	Calculation
Cloud compute costs	AWS/GCP/Azure charge by CPU-hour; efficient code uses fewer hours
Data center electricity	At \$0.10/kWh, a 10 MW data center costs ~\$8.7M/year in power
Cooling overhead	Every watt of compute requires 0.5-1.0 additional watts for cooling (PUE)
Carbon costs	Emerging carbon taxes (\$50-150/tonne CO ₂) add to operational expenses

Example: A Python service rewritten in an energy-efficient language (up to 75x improvement based on published benchmarks [6]) could reduce:

- Cloud compute costs by 50-90%
- Carbon footprint proportionally
- Cooling requirements in on-premise deployments

For CTOs evaluating technology choices, energy efficiency is no longer just an environmental consideration; it's a significant driver of operational expenditure.

5. Recommendations

For Individual Developers

1. **Profile for energy**, not just time—tools like PowerTOP, Intel Power Gadget, and perf can help [11]
2. **Choose efficient algorithms** and data structures
3. **Consider language energy characteristics** for energy-sensitive applications
4. **Avoid unnecessary abstraction** layers
5. **Test on target hardware** rather than assuming cloud resources are free

For Organizations

1. **Include energy metrics** in CI/CD pipelines
2. **Set energy budgets** for critical services
3. **Evaluate language choices** with energy in mind
4. **Invest in energy measurement** infrastructure
5. **Consider total cost of ownership**, including energy

For the Industry

1. **Develop energy measurement standards** for software
 2. **Create energy efficiency benchmarks** for languages and frameworks
 3. **Fund research** into energy-aware compilation
 4. **Integrate energy education** into computer science curricula
 5. **Reward energy efficiency** in hiring and promotion
-

6. Conclusion

The computing industry's energy footprint is substantial and growing. While hardware efficiency and renewable energy are important, they are insufficient solutions alone. Software, and the programming languages used to create it, must become more energy-efficient.

Sustainable computing is not about sacrifice or limitation. Energy-efficient software is often faster software. Energy-aware design encourages better architecture. Sustainable computing can be excellent computing.

The tools and techniques exist. The path forward is clear. What remains is the collective will to walk it.

References

- [1] International Energy Agency. "Data Centres and Data Transmission Networks." IEA, 2024. <https://www.iea.org/energy-system/buildings/data-centres-and-data-transmission-networks>
- [2] International Energy Agency. "Electricity 2024: Analysis and Forecast to 2026." IEA, January 2024. <https://www.iea.org/reports/electricity-2024>
- [3] Statista. "Number of smartphone subscriptions worldwide from 2016 to 2028." Statista, 2024. <https://www.statista.com/statistics/330695/number-of-smartphone-users-worldwide/>
- [4] Patterson, D., Gonzalez, J., Le, Q., et al. "Carbon Emissions and Large Neural Network Training." arXiv:2104.10350, 2021. <https://arxiv.org/abs/2104.10350>
- [5] Wirth, N. "A Plea for Lean Software." Computer, vol. 28, no. 2, pp. 64-68, 1995. <https://doi.org/10.1109/2.348001>

- [6] Pereira, R., Couto, M., Ribeiro, F., et al. “Energy Efficiency across Programming Languages.” Proceedings of the 10th ACM SIGPLAN International Conference on Software Language Engineering (SLE), 2017. <https://joule.openie.dev/research/energy-efficiency-languages>
- [7] Pinto, G., Castor, F., Liu, Y.D. “Mining Questions About Software Energy Consumption.” Proceedings of the 11th Working Conference on Mining Software Repositories (MSR), 2014. <https://doi.org/10.1145/2597073.2597110>
- [8] Mittal, S., Vetter, J.S. “A Survey of CPU-GPU Heterogeneous Computing Techniques.” ACM Computing Surveys, vol. 47, no. 4, 2015. <https://doi.org/10.1145/2788396>
- [9] Jouppi, N.P., Young, C., Patil, N., et al. “In-Datacenter Performance Analysis of a Tensor Processing Unit.” Proceedings of the 44th Annual International Symposium on Computer Architecture (ISCA), 2017. <https://doi.org/10.1145/3079856.3080246>
- [10] Schulte, E., Dorn, J., Harding, S., et al. “Post-compiler Software Optimization for Reducing Energy.” Proceedings of the 19th International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS), 2014. <https://doi.org/10.1145/2541940.2541980>
- [11] Khan, K.N., Hirki, M., Niemi, T., et al. “RAPL in Action: Experiences in Using RAPL for Power Measurements.” ACM Transactions on Modeling and Performance Evaluation of Computing Systems, vol. 3, no. 2, 2018. <https://doi.org/10.1145/3177754>
- [12] Drepper, U. “What Every Programmer Should Know About Memory.” Red Hat, Inc., 2007. <https://www.akkadia.org/drepper/cpumemory.pdf>



This work is licensed under CC BY 4.0 Open Access Research - Freely Shareable

Published by the Joule Project — joule-lang.org